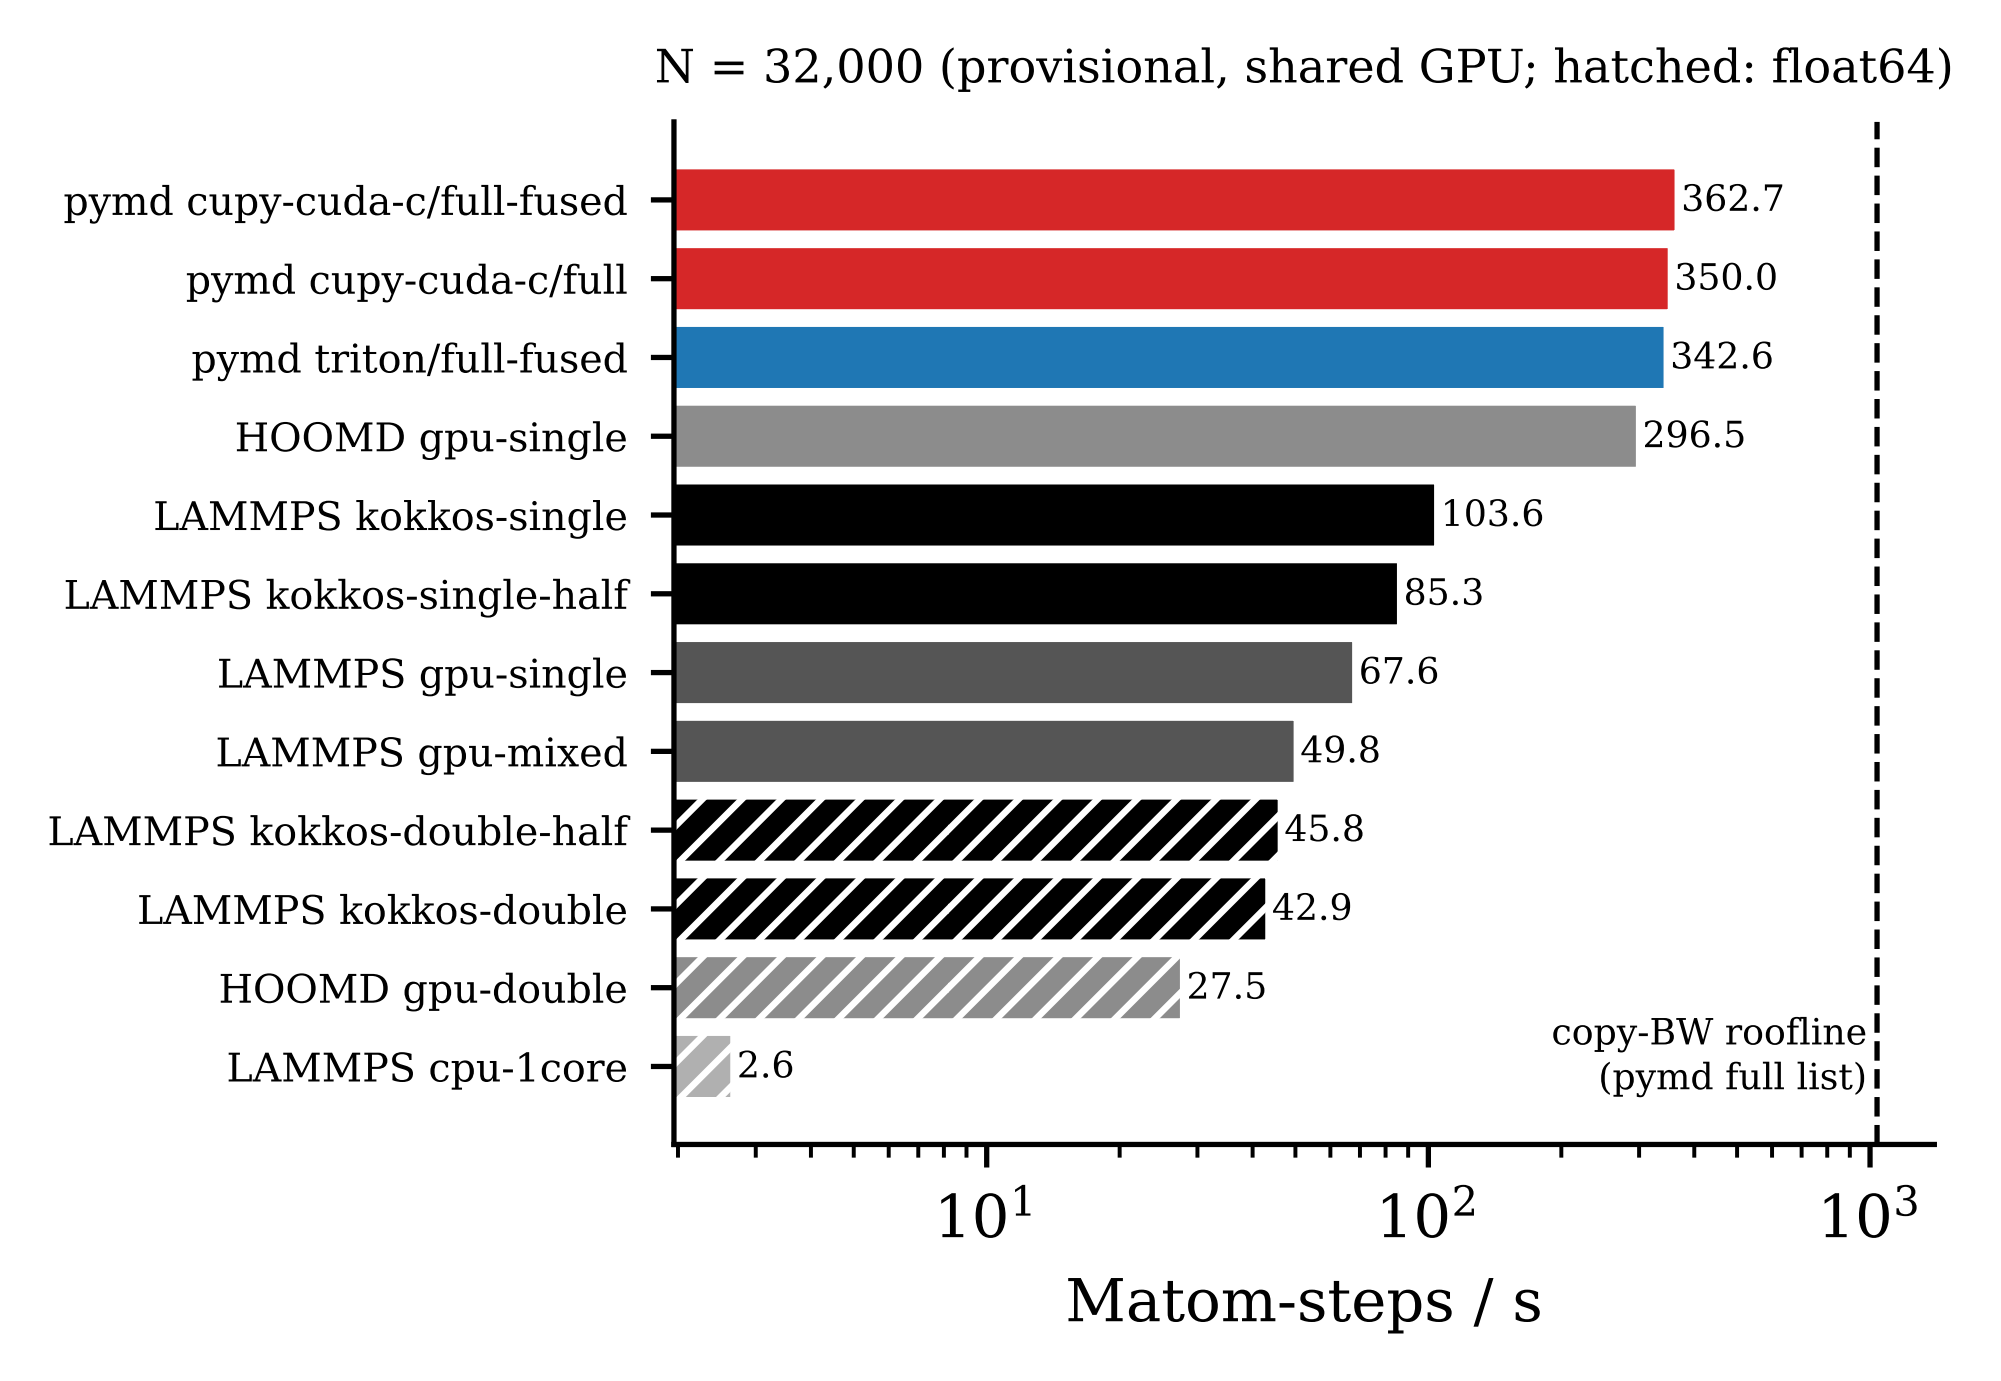
N = 32,000 (provisional, shared GPU; hatched: float64)
pymd cupy-cuda-c/full-fused
pymd cupy-cuda-c/full
pymd triton/full-fused
HOOMD gpu-single
LAMMPS kokkos-single
LAMMPS kokkos-single-half
LAMMPS gpu-single
LAMMPS gpu-mixed
LAMMPS kokkos-double-half
LAMMPS kokkos-double
HOOMD gpu-double
LAMMPS cpu-1core
362.7
350.0
342.6
296.5
103.6
85.3
67.6
49.8
45.8
42.9
27.5
2.6
copy-BW roofline
(pymd full list)
10^1
10^2
10^3
Matom-steps / s